Verified Inspection Desk

- Arabic / English municipal assistance
- The model proposes. The human decides. The database proves it.
- SDAIA Academy capstone | Track 1 | Synthetic data only

One workflow, end to end

- Verify the permit holder without revealing whether an unverified permit exists.
- Retrieve inspection policy over MCP; find actual district availability.
- Present an exact booking for human approval; commit only after revalidation.
- Demo fixture: FTR-30311 / health inspection / Al-Rawdah.

Keep authority outside the model

- CLI or authenticated HTTP client -> bounded agent loop
- Public policy MCP -> retrieved rules, no applicant database tools
- Local permit + slot + booking tools -> one persistent session database
- Human approval -> authorization recheck -> atomic write -> read-back

A prompt is not a permission system

- Four tools only; unknown capabilities are unavailable.
- Identity, current ownership and agent authorization checked in code.
- Denied approval stays denied for that turn; exact slot rechecked at commit.
- Bounded history/results, step and cost limits; redacted evidence.
- Offline verification: 102 checks; full evidence in results/.

Change one thing; measure honestly

- Suite: 16 original tasks; two configurations.
- A: generated final reply. B: verified booking/denial renderer. Safety unchanged.
- A: 1/2 completed tasks passed (50.0%); 2/16 runs complete
- B: 2/2 completed tasks passed (100.0%); 2/16 runs complete
- Paired synthetic response replay is disclosed; no repeated-trial reliability claim.
- INCOMPLETE: provider/quota or missing runs remain; do not present as a completed evaluation.

Failures are part of the deliverable

- A / health-approve-en: booking_id_missing
- Earlier smoke runs invented a reminder and answered in the wrong language.
- Verified outcome rendering prevents those booking-confirmation claims.
- Do not confuse one paired trial with pass^3 or production readiness.

Break a capability chain, not just a phrase

- Private records + untrusted ticket + outbound channel form the trifecta.
- A false audit override requested full identity fields to an attacker address.
- Sensitive outbox rows: 1 unsafe -> 0 defended
- The control is an isolated unsafe lab; the defended run uses our actual app.
- Only a local synthetic outbox is used. No real email. One attack is not comprehensive assurance.

Reproduce, inspect, question

- uv run rafeeq.py checks
- uv run rafeeq.py chat --session demo
- uv run rafeeq.py serve --host 127.0.0.1 --port 8793
- README + tasks + A/B CSVs + security evidence + slides.
- Limits: synthetic identity, isolated session databases, free-provider availability, heuristic PII masking.